

Aim - Develop a model-based RL algorithm, such as Monte Carlo Tree Search (MCTS), to solve a complex environment like Atari games.

Objective –

The objective of this experiment is to implement and train a Deep Q-Network (DQN) using PyTorch on the CartPole environment from OpenAI Gym. The goal is to learn an optimal policy to balance the pole on a cart by moving left or right.

Software Requirements –

- 1) Python (3.x recommended)
- 2) Jupyter Notebook or any Python IDE

Hardware Requirements - A machine with sufficient RAM and processing power for model training (8GB RAM recommended)

Prerequisites –

- 1) Basic understanding of Python programming
- 2) Familiarity with the concepts of Neural Networks

Libraries or Modules Used -

- 1) TensorFlow or PyTorch
- 2) NumPy
- 3) Matplotlib

Theory –

Methodology :

Environment: Utilize the CartPole environment provided by OpenAI Gym.

DQN Architecture: Design a neural network model to serve as the Q-network, which takes the state of the environment as input and outputs Q-values for each action.

Replay Buffer: Implement a replay buffer to store and sample experiences for training the DQN.

Training: Train the DQN using a combination of Q-learning and experience replay to optimize the Q-values.

Evaluation: Assess the performance of the trained DQN by measuring the average reward obtained over a specified number of episodes.

Hypothesis:

We hypothesize that the DQN will be able to learn an effective policy to balance the pole on the cart by maximizing the cumulative reward over episodes through reinforcement learning.

Experiment Setup:

Environment: CartPole-v1 environment from OpenAI Gym.

Neural Network: PyTorch-based DQN architecture.

Training Parameters: Learning rate, discount factor, exploration parameters.

Evaluation Metrics: Average reward per episode, stability of training.

Results and Analysis:

Analyze the performance of the trained DQN by examining the average reward obtained per episode and assessing the stability of training. Evaluate whether the DQN is able to learn an effective policy for balancing the pole on the cart.

Code -

```
import gym
import random
import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim

# Define Deep Q-Network (DQN) model
class DQN(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(DQN, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.fc2 = nn.Linear(hidden_size, hidden_size)
        self.fc3 = nn.Linear(hidden_size, output_size)

    def forward(self, x):
```

```
x = torch.relu(self.fc1(x))  
x = torch.relu(self.fc2(x))  
return self.fc3(x)
```

Define replay buffer

```
class ReplayBuffer:
```

```
    def __init__(self, capacity):
```

```
        self.capacity = capacity
```

```
        self.buffer = []
```

```
    def add(self, state, action, reward, next_state, done):
```

```
        experience = (state, action, reward, next_state, done)
```

```
        self.buffer.append(experience)
```

```
        if len(self.buffer) > self.capacity:
```

```
            self.buffer.pop(0)
```

```
    def sample(self, batch_size):
```

```
        return random.sample(self.buffer, batch_size)
```

Define Deep Q-Learning agent

```
class DQNAgent:
```

```
    def __init__(self, input_size, output_size, hidden_size=64, learning_rate=0.001, gamma=0.99,  
epsilon_start=1.0, epsilon_decay=0.995, epsilon_min=0.01, replay_buffer_capacity=10000,  
batch_size=64):
```

```
        self.input_size = input_size
```

```
        self.output_size = output_size
```

```
        self.hidden_size = hidden_size
```

```
        self.learning_rate = learning_rate
```

```
        self.gamma = gamma
```

```

self.epsilon = epsilon_start
self.epsilon_decay = epsilon_decay
self.epsilon_min = epsilon_min
self.replay_buffer = ReplayBuffer(replay_buffer_capacity)
self.batch_size = batch_size

self.q_network = DQN(input_size, hidden_size, output_size)
self.target_network = DQN(input_size, hidden_size, output_size)
self.target_network.load_state_dict(self.q_network.state_dict())
self.optimizer = optim.Adam(self.q_network.parameters(), lr=learning_rate)
self.loss_function = nn.MSELoss()

def epsilon_greedy_action(self, state):
    if np.random.rand() < self.epsilon:
        return np.random.choice(range(self.output_size))
    else:
        with torch.no_grad():
            q_values = self.q_network(torch.FloatTensor(state))
            return torch.argmax(q_values).item()

def train(self, state, action, reward, next_state, done):
    self.replay_buffer.add(state, action, reward, next_state, done)
    if len(self.replay_buffer.buffer) > self.batch_size:
        batch = self.replay_buffer.sample(self.batch_size)
        states, actions, rewards, next_states, dones = zip(*batch)

        states = torch.FloatTensor(states)
        actions = torch.LongTensor(actions)
        rewards = torch.FloatTensor(rewards)

```

```

next_states = torch.FloatTensor(next_states)
dones = torch.FloatTensor(dones)

q_values = self.q_network(states)
next_q_values = self.target_network(next_states).max(1)[0]
target_q_values = rewards + (1 - dones) * self.gamma * next_q_values

q_values = q_values.gather(1, actions.unsqueeze(1)).squeeze(1)

loss = self.loss_function(q_values, target_q_values.detach())

self.optimizer.zero_grad()
loss.backward()
self.optimizer.step()

self.epsilon = max(self.epsilon * self.epsilon_decay, self.epsilon_min)

def update_target_network(self):
    self.target_network.load_state_dict(self.q_network.state_dict())

# Initialize environment and agent
env = gym.make('CartPole-v1')
input_size = env.observation_space.shape[0]
output_size = env.action_space.n
agent = DQNAgent(input_size, output_size)

# Training
num_episodes = 1000
for episode in range(num_episodes):

```

```
state = env.reset()

total_reward = 0

done = False

while not done:

    action = agent.epsilon_greedy_action(state)

    next_state, reward, done, _ = env.step(action)

    agent.train(state, action, reward, next_state, done)

    state = next_state

    total_reward += reward

if episode % 100 == 0:

    agent.update_target_network()

    print(f"Episode {episode}, Total Reward: {total_reward}")

# Close the environment

env.close()
```

Conclusion : This code sets up a DQN agent using PyTorch and trains it on the CartPole environment. It includes the necessary components such as the DQN model, replay buffer, and training algorithm. The agent learns to balance the pole on the cart by maximizing cumulative rewards over episodes.